

Monitor class

Wrappers

Very frequently, you will want to extend the environment's functionality in some generic way. For example, an environment gives you some observations, but you want to accumulate them in some buffer and provide to the agent the N last observations, which is a common scenario for dynamic computer games, when one single frame is just not enough to get the full information about the game state. Another example is when you want to be able to crop or preprocess an image's pixels to make it more convenient for the agent to digest or if you want to normalize reward scores somehow. There are many such situations that have the same structure: you'd like to "wrap" the existing environment and add some extra logic doing something. Gym provides you with a convenient framework for these situations, called the Wrapper class. The class structure is shown in the following diagram. The Wrapper class inherits the Env class. Its constructor accepts the only argument: the instance of the Env class to be "wrapped." To add extra functionality, you need to redefine the methods you want to extend such as `step()` or `reset()`. The only requirement is to call the original method of the superclass.

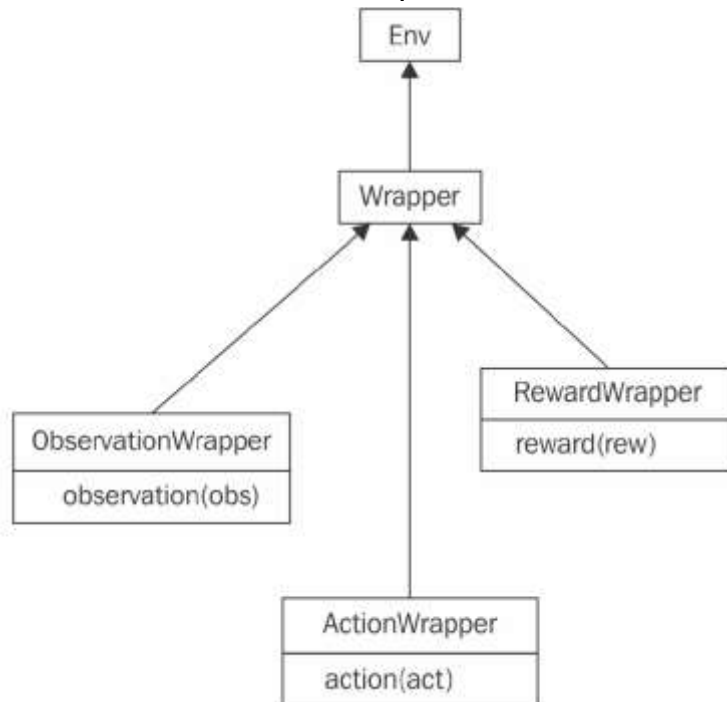


Figure : The hierarchy of Wrapper classes in Gym

To handle more specific requirements, such as a Wrapper class that wants to process only observations from the environment or only actions, there are subclasses of Wrapper that allow filtering of only a specific portion of information.

They are as follows: ObservationWrapper: You need to redefine observation(obs) method of the parent. The obs argument is an observation from the wrapped environment, and this method should return the observation that will be given to the agent.

RewardWrapper: This exposes the reward(rew) method, which could modify the reward value given to the agent. ActionWrapper: You need to override the action(act) method, which could tweak the action passed to the wrapped environment to the agent.

To make it slightly more practical, let's imagine a situation where we want to intervene in the stream of actions sent by the agent and, with a probability of 10%, replace the current action with a random one. It might look like an unwise thing to do, but this simple trick is one of the most practical and powerful methods to solving the "exploration/exploitation problem". By issuing the random actions, we make our agent explore the environment and from time to time drift away from the beaten track of its policy. This is an easy thing to do, using the ActionWrapper class (full example,

Deep Q Learning Agent/03_random_action_wrapper.py).

```
import gym
import random
class RandomActionWrapper(gym.ActionWrapper):
    def __init__(self, env, epsilon=0.1):
        super(RandomActionWrapper, self).__init__(env)
        self.epsilon = epsilon
```

Here we initialize our wrapper by calling a parent's __init__ method and saving epsilon (a probability of a random action):

```
def action(self, action):
    if random.random() < self.epsilon:
        print("Random!")
    return self.env.action_space.sample()
return action
```

This is a method that we need to override from a parent's class to tweak the agent's actions. Every time we roll the die and

with the probability of epsilon, we sample a random action from the action space and return it instead of the action the agent has sent to us. Note that using `action_space` and wrapper abstractions, we were able to write abstract code, which will work with *any* environment from the Gym. Additionally, we print the message every time we replace the action, just to verify that our wrapper is working. In the production code, of course, this won't be necessary:

```
if __name__ == "__main__":  
    env = RandomActionWrapper(gym.make("CartPole-v0"))
```

Now it's time to apply our wrapper. We will create a normal `CartPole` environment and pass it to our wrapper constructor. From here on, we use our wrapper as a normal `Env` instance, instead of the original `CartPole`. As the `Wrapper` class inherits the `Env` class and exposes the same interface, we can nest our wrappers in any combination we want. This is a powerful, elegant, and generic solution:

```
obs = env.reset()  
total_reward = 0.0  
while True:  
    obs, reward, done, _ = env.step(0)  
    total_reward += reward  
    if done:  
        break  
    print("Reward got: %.2f" % total_reward)
```

Here is almost the same code, except that every time we issue the same action: 0.

Our agent is dull and always does the same thing. By running the code, you should see that the wrapper is indeed working:

```
$ python 03_random_actionwrapper.py  
WARN: gym.spaces.Box autodetected dtype as <class 'numpy.float32'>.  
Please provide explicit dtype.  
Random!  
Random!  
Random!  
Random!  
Reward got: 12.00
```

If you want, you can play with the epsilon parameter on the wrapper's creation and verify that randomness improves the agent's score on average. We should move on and look at another interesting gem hidden inside Gym: Monitor.

Monitor

Another class you should be aware of is Monitor. It is implemented like Wrapper and can write information about your agent's performance in a file with an optional video recording of your agent in action.

Just to give you an idea of how the Gym web interface looked, here is the CartPole environment leaderboard:

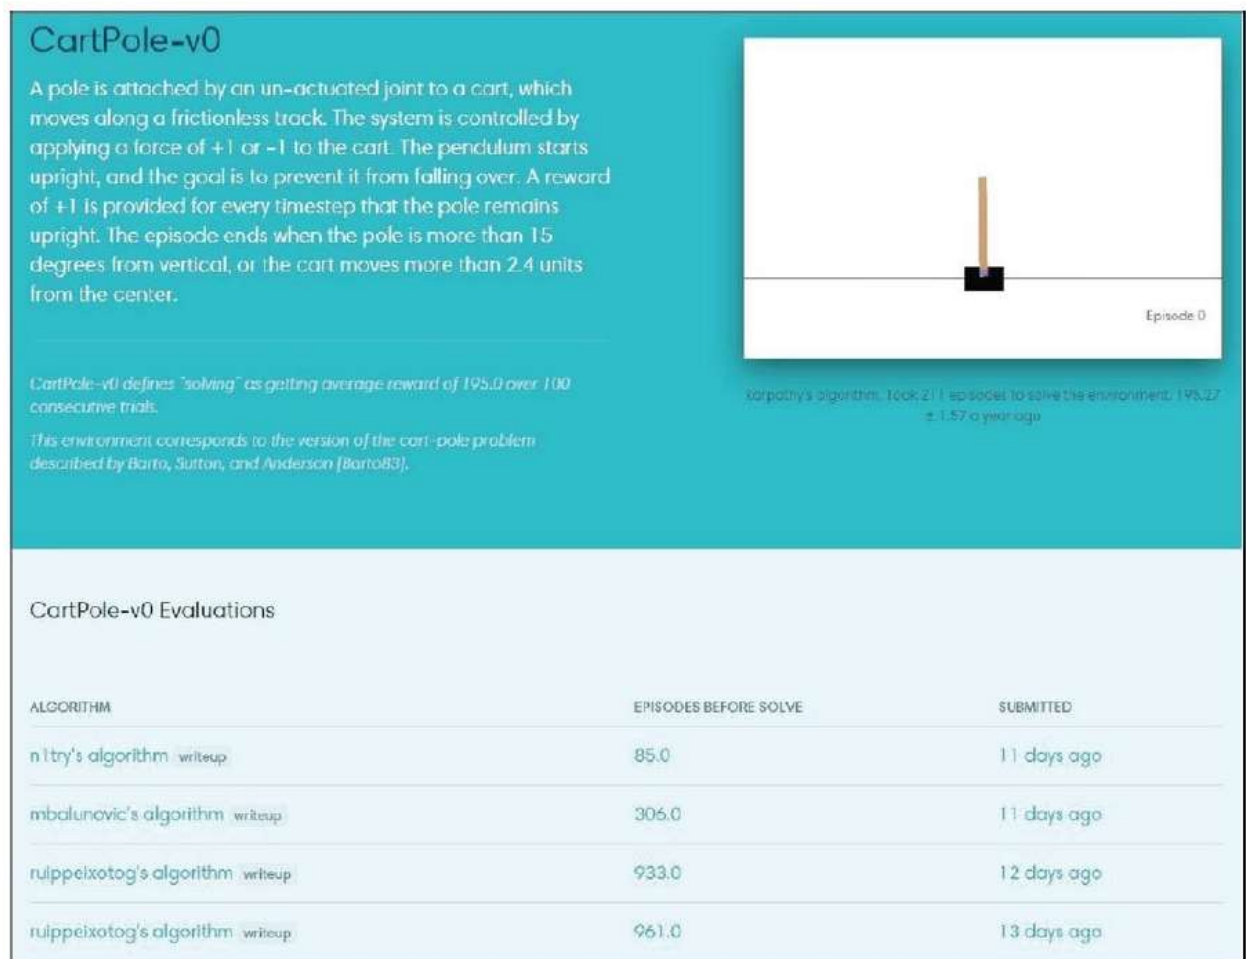


Figure : OpenAI Gym web interface with CartPole submissions

Despite this, Monitor is still useful, as you can take a look at your agent's life inside the environment. So, here is how we add Monitor to our random CartPole agent, which is the only difference (the entire code is in

Deep Q Learning Agent /04_cartpole_random_monitor.py) :

```
if __name__ == "__main__":  
    env = gym.make("CartPole-v0")  
    env = gym.wrappers.Monitor(env, "recording")
```

The second argument that we pass to Monitor is the name of the directory it will write the results to. This directory shouldn't exist, otherwise your program will fail with an exception (to overcome this, you could either remove the existing directory or pass the `force=True` argument to the Monitor class' constructor).

The Monitor class requires the FFmpeg utility to be present on the system, which is used to convert captured observations into an output video file. This utility must be available, otherwise Monitor will raise an exception. The easiest way to install FFmpeg is using your system's package manager, which is OS distribution specific.

To start this example, one of these three extra prerequisites should be met:

- The code should be run in an X11 session with the OpenGL extension (GLX).

- The code should be started in an xvfb virtual display.

- You can use X11 forwarding in ssh connection.

The cause of this is video recording, which is done by taking screenshots of the window drawn by the environment. Some of the environment uses OpenGL to draw its picture, so the graphical mode with OpenGL needs to be present. This could be a problem for a virtual machine in the cloud, which physically doesn't have a monitor and graphical interface running. To overcome this, there is a special "virtual" graphical display, called **Xvfb (X11 virtual framebuffer)**, which basically starts a virtual graphical display on the server and forces the program to draw inside it. This would be enough to make Monitor happily create the desired videos.

To start your program in the `xvfb` environment, you need to have it installed on your machine (it usually requires installing the `xvfb` package) and run the special script,

```
xvfb-run:
```

```
$ xvfb-run -s "-screen 0 640x480x24" python
```

```
04_cartpole_random_monitor.py
```

```
[2017-09-22 12:22:23,446] Making new env: CartPole-v0
```

```
[2017-09-22 12:22:23,451] Creating monitor directory recording
```

```
[2017-09-22 12:22:23,570] Starting new video recorder writing to  
recording/openaigym.video.0.31179.video000000.mp4
```

```
Episode done in 14 steps, total reward 14.00
```

```
[2017-09-22 12:22:26,290] Finished writing results. You can upload  
them to the scoreboard via gym.upload('recording')
```

As you may see from the preceding log, the video has been written successfully, so you can peek inside one of your agent's sections by playing it. Another way to record your agent's actions is to use ssh X11 forwarding, which uses the ssh ability to tunnel X11 communications between the X11 client (Python code which wants to display some graphical information) and X11 server (software which knows how to display this information and has access to your physical display). In X11 architecture, the client and the server are separated and can work on different machines.

To use this approach, you need the following:

1. An X11 server running on your local machine. Linux comes with X11 server as a standard component (all desktop environments are using X11).

On a Windows machine, you can set up third-party X11 implementations such as open source `VcXsrv` (available in <https://sourceforge.net/projects/vcxsrv/>).

2. The ability to log in to your remote machine via ssh, passing the `-X` command-line option: `ssh -X servername`. This enables X11 tunneling and allows all processes started in this session to use your local display for graphics output.

Then you can start a program that uses the `Monitor` class and it will display the agent's actions, capturing the images into a video file.